

## **Assignment 1 Guidelines**

### **Machine Learning and Artificial Intelligence in Finance**

***FINC 5322***

***Fall 2025***

**The American University in Cairo**

**Onsi Sawiris School of Business**

**Heikal Department of Management**

### ***Customer Churn Prediction – Assignment 1***

In business, customer churn (or attrition) refers to when customers stop doing business with a company.

For example, in a subscription-based telecom company, a customer is considered churned if they cancel their subscription.

Predicting churn is critical because:

- Retaining an existing customer is often cheaper than acquiring a new one.
- Companies can take action (targeted campaigns, discounts, improved services) to prevent churn.
- Understanding who is likely to churn helps businesses allocate resources more effectively.

In this assignment, you will work with a telecom churn dataset. Your goal is not to build the most accurate model but to practice the full data preparation pipeline: exploring, cleaning, engineering features, and running a simple model.

---

### ***Accessing the Assignment***

- All Assignments are hosted on our course Code Repository (GitHub):  
<https://github.com/kelkess43/AUC-Material>
- The Assignment will also be shared through Canvas modules for easy access

---

### ***Structure of the Assignment***

Your notebook is divided into five sections.

You must complete the TODOs in the notebook and keep all work within the. ipynb file.

#### ***Section 1: Exploratory Data Analysis (EDA)***

- Count the number of unique customers.
- Check for missing values.
- Visualize numeric distributions (tenure, charges)
- Visualize churn distribution (how many customers stayed vs churned).

### *Section 2: Data Cleaning*

- Convert TotalCharges to numeric and handle missing values.
- Remove duplicate rows.
- Reflect briefly: Why did you choose your method for handling missing values?

### *Section 3: Feature Engineering*

- Create a tenure\_group feature (e.g., 0–12 months, 13–24 months, etc.).
- Encode categorical variables into numeric form.

### *Section 4: Simple Modelling*

- Perform train/test split (70/30).
- Train a Logistic Regression model.
- Evaluate with accuracy, precision, and recall.

### *Section 5: Research & Reflection*

- Explain why accuracy is not enough in churn prediction.
- Research and explain in your own words what precision, recall, and the classification report mean.

---

### ***Deliverables and Submission***

- Submit one Jupyter Notebook (.ipynb) file with all your work.
- Make sure all cells are executed, and outputs are visible.
- Add short comments/markdown where required
- Deadline: Tuesday, 7th October at 11:59 PM
- Submission: Upload your completed. ipynb notebook on Canvas

---

### ***Grading***

- EDA (20 points)
- Cleaning (20 points)
- Feature Engineering (20 points)
- Modelling (25 points)
- Research & Reflection (15 points)